

# Topics in R: Tidyverse for Data Analysis

UCI Data Science Initiative

Dustin Pluta

2018-05-04

# Overview

## AM

- Intro to the Tidyverse
- Importing and wrangling data with `readr` and `tidyr`.
- Exploration and visualization with `ggplot2` and `dplyr`.

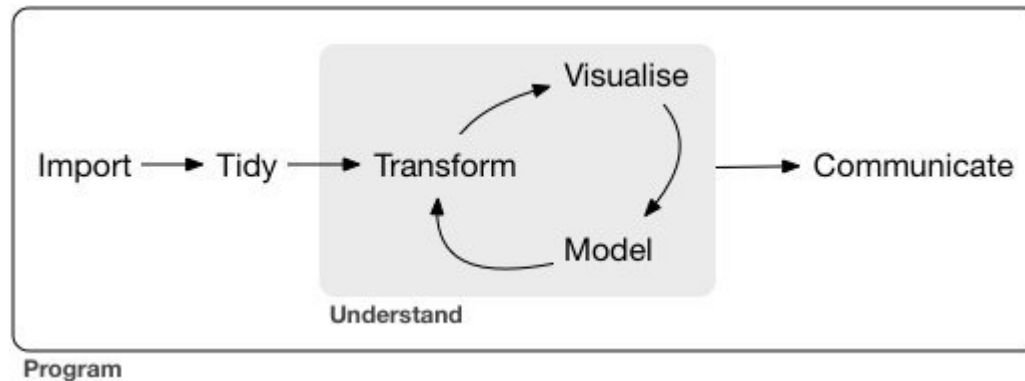
## Lunch

## PM

- Tidyverse in application: Data Analysis of Coronary Heart Disease Data
- Writing reports with R Markdown.

# Intro to the Tidyverse

## Tidy Analysis Pipeline



- *tidyverse* philosophy: collection of small, simple functions that each do one thing well
- Written by Hadley Wickham, Chief Scientist for R Studio, who also developed:
  - `ggplot2`
  - `reshape2`
  - `tidyr`
  - many others

# Intro to the Tidyverse

## Tidy Data

country	year	cases	population
Afghanistan	1999	37737	19987071
Afghanistan	2000	3666	20395360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	214258	1272015272
China	2000	216766	1280425583

variables

country	year	cases	population
Afghanistan	1999	37737	19987071
Afghanistan	2000	3666	20395360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	214258	1272015272
China	2000	216766	1280425583


observations

country	year	cases	population
Afghanistan	99	37737	19987071
Afghanistan	00	3666	20395360
Brazil	99	37737	172006362
Brazil	00	80488	174504898
China	99	214258	1272015272
China	00	216766	1280425583

values

# Intro to the Tidyverse

Book: R for Data Science



<http://r4ds.had.co.nz/>

# Intro to the Tidyverse

## Packages

**readr:** import and export data

**tidyr:** wrangle and clean data

**dplyr:** slice, subset, transform, and summarize data

**ggplot2:** visualization

**purrr:** vectorized and parallel operations

**RMarkdown:** preparing and presenting results

---

# Intro to Tidyverse: Getting Started

## Install Packages

```
install.packages("tidyverse")  
install.packages("rmarkdown")
```

# Load Tidyverse

```
library(tidyverse)
```



# Load IMDB Data

```
imdb <- read_csv("dat/movie_metadata.csv")
```

```
## Parsed with column specification:
## cols(
##   .default = col_integer(),
##   color = col_character(),
##   director_name = col_character(),
##   actor_2_name = col_character(),
##   genres = col_character(),
##   actor_1_name = col_character(),
##   movie_title = col_character(),
##   actor_3_name = col_character(),
##   plot_keywords = col_character(),
##   movie_imdb_link = col_character(),
##   language = col_character(),
##   country = col_character(),
##   content_rating = col_character(),
##   imdb_score = col_double(),
##   aspect_ratio = col_double()
## )
```

```
## See spec(...) for full column specifications.
```

# IMDB Data

```
head(imdb)
```

```
## # A tibble: 6 x 28
##   color director_name num_critic_for_... duration director_facebo...
##   <chr> <chr>                <int>      <int>                <int>
## 1 Color James Cameron          723        178                  0
## 2 Color Gore Verbins...        302        169                 563
## 3 Color Sam Mendes             602        148                  0
## 4 Color Christopher ...        813        164                22000
## 5 <NA> Doug Walker              NA         NA                  131
## 6 Color Andrew Stant...        462        132                 475
## # ... with 23 more variables: actor_3_facebook_likes <int>,
## #   actor_2_name <chr>, actor_1_facebook_likes <int>, gross <int>,
## #   genres <chr>, actor_1_name <chr>, movie_title <chr>,
## #   num_voted_users <int>, cast_total_facebook_likes <int>,
## #   actor_3_name <chr>, facenumber_in_poster <int>, plot_keywords <chr>,
## #   movie_imdb_link <chr>, num_user_for_reviews <int>, language <chr>,
## #   country <chr>, content_rating <chr>, budget <int>, title_year <int>,
## #   actor_2_facebook_likes <int>, imdb_score <dbl>, aspect_ratio <dbl>,
## #   movie_facebook_likes <int>
```

# IMDB Data

```
colnames(imdb)
```

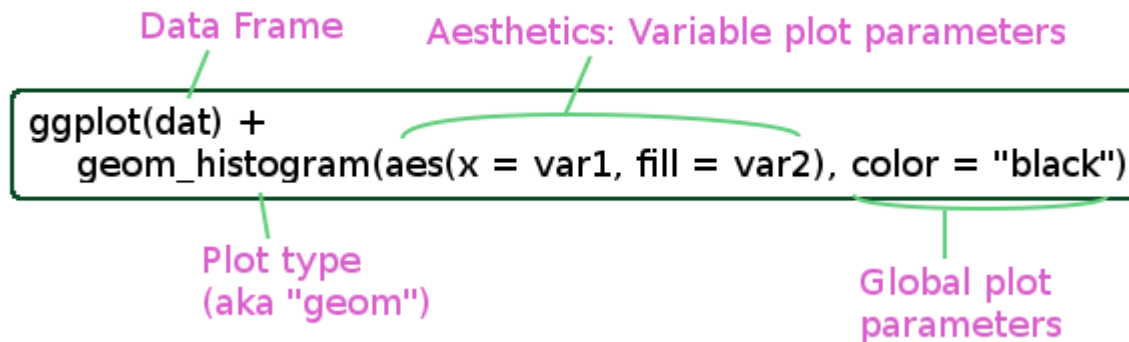
```
## [1] "color" "director_name"
## [3] "num_critic_for_reviews" "duration"
## [5] "director_facebook_likes" "actor_3_facebook_likes"
## [7] "actor_2_name" "actor_1_facebook_likes"
## [9] "gross" "genres"
## [11] "actor_1_name" "movie_title"
## [13] "num_voted_users" "cast_total_facebook_likes"
## [15] "actor_3_name" "facenumber_in_poster"
## [17] "plot_keywords" "movie_imdb_link"
## [19] "num_user_for_reviews" "language"
## [21] "country" "content_rating"
## [23] "budget" "title_year"
## [25] "actor_2_facebook_likes" "imdb_score"
## [27] "aspect_ratio" "movie_facebook_likes"
```

# Visualization with ggplot2

- ggplot2 is a plotting package that is a nice and more modern alternative to R base plots
- Based on the idea of a *grammar of graphics*...
  - Think of a **plot like a sentence**...
  - **noun**: the plot data
  - **verbs**: the plot types
  - **adverbs**: the plot characteristics

# Visualization with `ggplot2`

# Visualization with ggplot2



# Visualization with `ggplot2`

## List of geoms


- `geom_histogram`
- `geom_density`
- `geom_point`
- `geom_line`
- `geom_boxplot`
- ... many others

# Visualization with ggplot2

```
library(ggplot2)
```

```
ggplot(imdb) +  
  geom_histogram(aes(x = imdb_score))
```

## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.





# Visualization with ggplot2

## Some Aesthetics

*x*: horizontal position

*y*: vertical position

*alpha*: transparency

*color*: border color

*fill*: interior color

*group*: grouping variable

*linetype*

*size*

# Visualization with ggplot2

**Different geoms have different aesthetics**


**Refer to the documentation to see which aesthetics are supported for a geom**

```
?geom_histogram
```

# Visualization with ggplot2


```
imdb$recent <- imdb$title_year > 2000  
ggplot(imdb) +  
  geom_histogram(aes(x = imdb_score, fill = recent))
```

## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.




# Visualization with ggplot2

```
imdb$recent <- imdb$title_year > 2000  
ggplot(imdb) +  
  geom_histogram(aes(x = imdb_score, fill = recent),  
    color = "black", bins = 30)
```




# Visualization with ggplot2

```
ggplot(imdb) +  
  geom_histogram(aes(x = imdb_score, fill = recent),  
                 color = "black", bins = 30) +  
  xlab("IMDB Score") +  
  ylab("Count") +  
  ggtitle("Comparing IMDB Scores for Recent vs Old Movies")
```




# Visualization with ggplot2

```
ggplot(imdb) +  
  geom_histogram(aes(x = imdb_score, fill = recent),  
                 color = "black", bins = 30, alpha = 0.4) +  
  xlab("IMDB Score") +  
  ylab("Count") +  
  ggtitle("Comparing IMDB Scores for Recent vs Old Movies")
```




# Visualization with ggplot2

```
ggplot(imdb) +  
  geom_density(aes(x = imdb_score, fill = recent),  
               color = "black", alpha = 0.4) +  
  xlab("IMDB Score") +  
  ylab("Density") +  
  ggtitle("Comparing IMDB Scores for Recent vs Old Movies")
```



# Visualization with ggplot2


```
ggplot(imdb) +  
  geom_point(aes(x = budget, y = imdb_score))
```






# Visualization with ggplot2

```
ggplot(imdb) +  
  geom_point(aes(x = log(budget), y = imdb_score))
```




# Visualization with ggplot2

```
ggplot(imdb) +  
  geom_point(aes(x = log(budget), y = imdb_score), color = "green4")
```



# Visualization with ggplot2

```
ggplot(imdb) +  
  geom_point(aes(x = log(budget), y = imdb_score),  
             pch = 21, fill = "green4",  
             color = "black", alpha = 0.8)
```




# Visualization Exercises

1. Plot a histogram of budget and compare it to a histogram of `log(budget)`.
2. Add some color, change the title and axis labels for the `log(budget)` histogram.
3. Make a new variable `recent` to indicate if a movie is more recent than 2000 using `imdb$recent <- imdb$title_year > 2000`, then plot a histogram of `log(budget)` grouped by `recent`.
4. Create a scatterplot of `imdb_score` by `log(budget)` and colored by `recent`.
5. Create a boxplot of `imdb_score` grouped by `recent`, using `geom_boxplot`.

# Visualization Exercise # 1

```
ggplot(imdb) +  
  geom_histogram(aes(x = budget))
```


## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



# Visualization Exercise # 1


```
ggplot(imdb) +  
  geom_histogram(aes(x = log(budget)))
```

## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



# Visualization Exercise # 2


```
ggplot(imdb) +  
  geom_histogram(aes(x = log(budget)),  
                 color = "orange", fill = "purple", bins = 30) +  
  ggtitle("Exercise 2") +  
  xlab("The Log of the Budget") +  
  ylab("Frequency Count")
```



# Visualization Exercise # 3

```
imdb$recent <- imdb$title_year > 2000  
ggplot(imdb) +  
  geom_histogram(aes(x = log(budget), fill = recent),  
                 color = "black")
```


## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.






# Visualization Exercise # 4

```
ggplot(imdb) +  
  geom_point(aes(x = log(budget), y = imdb_score, fill = recent),  
             pch = 21, color = "black")
```




# Visualization Exercises # 5

```
ggplot(imdb) +  
  geom_boxplot(aes(y = imdb_score, x = recent, fill = recent)) +  
  coord_flip()
```




```
ggplot(imdb) +  
  geom_qq(aes(sample = imdb_score)) +  
  geom_qq_line(aes(sample = imdb_score))
```

```
ggplot(imdb) +  
  geom_point(aes(x = log(budget), y = imdb_score, fill = recent),  
             pch = 21, color = "black") +  
  geom_smooth(aes(x = log(budget), y = imdb_score, color = recent),  
             method = "lm")
```



# Visualization with ggplot2

```
ggplot(imdb[imdb$recent, ]) +  
  geom_boxplot(aes(y = imdb_score, fill = factor(title_year),  
    x = factor(title_year)))
```



# Data Wrangling with `dplyr`

# Data Wrangling with `dplyr`

- `dplyr` is a package designed for easy and efficient data manipulation

## Key Functions

- `filter`: select subset of rows (observations)
- `select`: select subset of columns (variables)
- `mutate`: transform variables in a data set
- `arrange`: reorder rows
- `summarise`: collapses a data frame into a single row
- `group_by`

# Data Wrangling with dplyr

- Let's try some dplyr functions with the iris data set:

```
# Print species means of Sepal Width
data(iris)
iris <- filter(iris, Species!="setosa")
iris <- select(iris, c(Sepal.Width, Species))
iris <- group_by(iris, Species)
species_means <- summarise(iris, mean(Sepal.Width))
print(species_means)
```

```
## # A tibble: 2 x 2
##   Species      `mean(Sepal.Width)`
##   <fct>                <dbl>
## 1 versicolor          2.77
## 2 virginica           2.97
```

# Data Wrangling with `dplyr`

## Introducing the Pipe: `%>%`

- `dplyr` (and much of the *tidyverse*) is designed around the use of the pipe operator `%>%`
- The pipe operator `%>%` allows you to chain operations on a data set together without having to create specific intermediate objects
- When using `%>%`, the first argument to a function is taken as the output of the previous step in the chain



# Data Wrangling with dplyr

- For example, the following is equivalent to the previous code:

```
# Prints species means, does not save anything  
# Original data.frame iris is unaffected  
data(iris)  
iris %>% filter(Species!="setosa") %>%  
  select(c(Sepal.Width, Species)) %>%  
  group_by(Species) %>%  
  summarise(mean(Sepal.Width))
```

```
## # A tibble: 2 x 2  
##   Species    `mean(Sepal.Width)`  
##   <fct>          <dbl>  
## 1 versicolor      2.77  
## 2 virginica       2.97
```

# Data Wrangling with dplyr

```
# To save the results instead
species_means <- iris %>%
  filter(Species!="setosa") %>%
  select(c(Sepal.Width, Species)) %>%
  group_by(Species) %>%
  summarise(mean(Sepal.Width))
```

```
species_means
```

```
## # A tibble: 2 x 2
##   Species      `mean(Sepal.Width)`
##   <fct>                <dbl>
## 1 versicolor          2.77
## 2 virginica           2.97
```

# dp1yr Exercises

1. Use `dplyr` to calculate the mean Sepal Width of the virginica species.
2. `summarise` can summarise multiple variables simultaneously, applying a (possibly different) function to each variable.  
Adapt the code below to find the minimum, median, maximum, and standard deviation of the Sepal.Width for the virginica species.
3. `group_by()` makes `summarise` even more useful by allowing you to summarise values across groups of a category simultaneously.  
Using `group_by`, adapt your code from the previous problem to produce the summary values for each species.

**Modify this code for problems 2 and 3:**

```
data(iris)
iris %>% summarise(mean_sepal_width = mean(Sepal.Width),
                   min_sepal_width = min(Sepal.Width))
```

# dplyr Exercises

## *Solution*

- (1) Use dplyr to calculate the mean Sepal Width of the virginica species.

```
data(iris)

iris %>%
  filter(Species == "virginica") %>%
  summarise(mean_sepal_width = mean(Sepal.Width))
```

```
##    mean_sepal_width
## 1             2.974
```

# dp1yr Exercises

## *Solution*

- (2) summarise can summarise multiple variables simultaneously, applying a (possibly different) function to each variable.  
Adapt the code below to find the minimum, median, maximum, and standard deviation of the Sepal.Width for the virginica species.

```
data(iris)
iris %>%
  filter(Species == "virginica") %>%
  summarise(min_sepal_width = min(Sepal.Width),
            med = median(Sepal.Width), maximum = max(Sepal.Width),
            stdev = sd(Sepal.Width))
```

```
##   min_sepal_width med maximum    stdev
## 1             2.2   3      3.8 0.3224966
```

# dpLyr Exercises

## *Solution*

- (3) `group_by()` makes `summarise` even more useful by allowing you to summarise values across groups of a category simultaneously.  
Using `group_by`, adapt your code from the previous problem to produce the summary values for each species.

```
data(iris)
iris %>%
  group_by(Species) %>%
  summarise(min_sepal_width = min(Sepal.Width),
            med = median(Sepal.Width), maximum = max(Sepal.Width),
            stdev = sd(Sepal.Width))
```

```
## # A tibble: 3 x 5
##   Species    min_sepal_width    med maximum    stdev
##   <fct>          <dbl> <dbl>    <dbl> <dbl>
## 1 setosa          2.3    3.4      4.4  0.379
## 2 versicolor      2     2.8      3.4  0.314
## 3 virginica       2.2    3       3.8  0.322
```

# Back to the Movies

# Exploring the IMDB Data

**How many movies for each actor in the dataset?**

```
imdb %>%  
  group_by(actor_1_name) %>%  
  summarize(n())
```



# Exploring the IMDB Data

**How many movies for each actor in the dataset?**

```
imdb %>%  
  group_by(actor_1_name) %>%  
  summarize(n())
```

```
## # A tibble: 2,098 x 2  
##   actor_1_name      `n()`  
##   <chr>           <int>  
## 1 50 Cent                1  
## 2 Aaliyah                1  
## 3 Aasif Mandvi           1  
## 4 Abbie Cornish          3  
## 5 Abhishek Bachchan      2  
## 6 Abigail Evans          1  
## 7 Abigail Spencer        1  
## 8 Adam Arkin             2  
## 9 Adam Baldwin          4  
## 10 Adam Garcia           3  
## # ... with 2,088 more rows
```

# Exploring the IMDB Data

**How many movies for each actor in the dataset?**

**Arranged by decreasing number of movies**

```
imdb %>%  
  group_by(actor_1_name) %>%  
  summarize(n_movies = n()) %>%  
  arrange(desc(n_movies))
```

```
## # A tibble: 2,098 x 2  
##   actor_1_name      n_movies  
##   <chr>          <int>  
## 1 Robert De Niro      49  
## 2 Johnny Depp        40  
## 3 Nicolas Cage       32  
## 4 J.K. Simmons       31  
## 5 Bruce Willis       30  
## 6 Denzel Washington  30  
## 7 Matt Damon        30  
## 8 Liam Neeson       29  
## 9 Harrison Ford     27  
## 10 Robin Williams   27
```

# Exploring the IMDB Data

**How many movies for each actor in the dataset?**

**Arranged by decreasing mean IMDB score**

```
imdb %>%  
  group_by(actor_1_name) %>%  
  summarize(mean_imdb_score = mean(imdb_score)) %>%  
  arrange(desc(mean_imdb_score))
```

```
## # A tibble: 2,098 x 2  
##   actor_1_name      mean_imdb_score  
##   <chr>              <dbl>  
## 1 Krystyna Janda      9.1  
## 2 Jack Warden         8.9  
## 3 Rob McElhenney      8.8  
## 4 Abigail Evans       8.7  
## 5 Elina Abai Kyzy     8.7  
## 6 Jackie Gleason      8.7  
## 7 Kimberley Crossman  8.7  
## 8 Maria Pia Calzone   8.7  
## 9 Takashi Shimura     8.7  
## 10 Bunta Sugawara     8.6
```

# Exploring the IMDB Data

**Considering actors with more than 5 movies, list top 10 actors with highest mean IMDB scores in decreasing order.**

```
imdb %>%  
  group_by(actor_1_name) %>%  
  summarize(mean_imdb_score = mean(imdb_score), n_movies = n()) %>%  
  filter(n_movies > 5) %>%  
  top_n(10, mean_imdb_score) %>%  
  arrange(desc(mean_imdb_score))
```

```
## # A tibble: 10 x 3  
##   actor_1_name      mean_imdb_score n_movies  
##   <chr>            <dbl>      <int>  
## 1 Leonardo DiCaprio      7.50        21  
## 2 Tom Hanks              7.42        24  
## 3 Clint Eastwood         7.34        16  
## 4 Tom Hardy              7.31        11  
## 5 Alan Rickman           7.29         8  
## 6 Benedict Cumberbatch  7.29         7  
## 7 Philip Seymour Hoffman  7.24        20  
## 8 Toby Jones             7.22         6  
## 9 Minnie Driver          7.21         7
```

# Exploring the IMDB Data

**How many movie entries does Harrison Ford have?**

```
actor <- "Harrison Ford"  
filter(imdb, actor_1_name == actor) %>% nrow
```

```
## [1] 27
```

# Exploring the IMDB Data

## Harrison Ford IMDB Scores


```
imdb %>%  
  filter(actor_1_name == actor) %>%  
  select(imdb_score)
```

```
## # A tibble: 27 x 1  
##   imdb_score  
##   <dbl>  
## 1      6.2  
## 2      6.7  
## 3      6.6  
## 4      6.1  
## 5      6.4  
## 6      5.7  
## 7      5.3  
## 8      5.1  
## 9      6.9  
## 10     6.3  
## # ... with 17 more rows
```

# Exploring the IMDB Data

## Histogram of Harrison Ford IMDB Scores

```
ggplot(imdb %>% filter(actor_1_name == actor)) +  
  geom_histogram(aes(x = imdb_score),  
                 fill = "darkred", color = "black", bins = 15)
```



# Exploring the IMDB Data

## Harrison By Genre

```
imdb %>%  
  filter(actor_1_name %in% actor) %>%  
  group_by(genres) %>%  
  summarize(mean_score = mean(imdb_score), n_movies = n())
```

```
## # A tibble: 22 x 3  
##   genres                                mean_score n_movies  
##   <chr>                                <dbl>      <int>  
## 1 Action|Adventure                      8.05         2  
## 2 Action|Adventure|Comedy|Romance       5.7          1  
## 3 Action|Adventure|Crime|Drama|Mystery|Thriller 7.8          1  
## 4 Action|Adventure|Drama|Thriller       6.4          1  
## 5 Action|Adventure|Fantasy             7.25         2  
## 6 Action|Adventure|Fantasy|Sci-Fi      8.63         3  
## 7 Action|Comedy|Crime|Thriller          5.3          1  
## 8 Action|Crime|Drama|Thriller           6.5          2  
## 9 Action|Drama|War                     6.3          1  
## 10 Action|Sci-Fi                       6.7          1  
## # ... with 12 more rows
```



# Exploring the IMDB Data

## Harrison's Action Movies

```
imdb %>%  
  filter(actor_1_name %in% actor) %>%  
  transmute(action = str_detect(genres, "Action"))
```

```
## # A tibble: 27 x 1  
##   action  
##   <lgl>  
## 1 TRUE  
## 2 TRUE  
## 3 FALSE  
## 4 TRUE  
## 5 TRUE  
## 6 TRUE  
## 7 TRUE  
## 8 FALSE  
## 9 TRUE  
## 10 FALSE  
## # ... with 17 more rows
```

# Exploring the IMDB Data

## Harrison's Action Movies

```
imdb %>%  
  filter(actor_1_name %in% actor) %>%  
  mutate(action = str_detect(genres, "Action")) %>%  
  group_by(action) %>%  
  summarize(mean(imdb_score))
```

```
## # A tibble: 2 x 2  
##   action `mean(imdb_score)`  
##   <lgl>         <dbl>  
## 1 FALSE         7.15  
## 2 TRUE          7.16
```

# Tidyverse in Application

# Tidyverse for Data Analysis

Need additional packages:

- boot
- factoextra
- FactoMineR
- qwraps2
- broom

## Goals

- Practice using the tidyverse for analysis of a real world data set
- Will use data from a study of Coronary Heart Disease (CHD) in older adults
- Interest in determining factors associated with higher incidence rate of CHD

# Coronary Heart Disease Data

- Covariates collected include:
  - `choltot`: Total cholesterol
  - `incchd`: Binary indicator of CHD
  - `hdl`: High-density lipoproteins
  - `alcoh`: Level of alcohol consumption
  - `trig`: Triglyceride level
  - `age`: Age in years
  - `diabetes`: Indicator of diabetes (3 levels: 1/2/3 for No, Type I, Type II)
  - `diabp`: Diastolic blood pressure
  - `sysbp`: Systolic blood pressure
  - `height`
  - `weight`
  - `smoke`: Binary indicator
  - `racebw`: Binary indicator of race white/black (0/1)
  - `gender`: Binary indicator of male/female (0/1)

# CHD Data

```
dat <- read_csv("dat/chddata_dsi.csv")
```

```
## Parsed with column specification:
```

```
## cols(  
##   X1 = col_integer(),  
##   idno = col_integer(),  
##   choltot = col_double(),  
##   incchd = col_integer(),  
##   hdl = col_integer(),  
##   alcohol = col_double(),  
##   bmi = col_double(),  
##   trig = col_integer(),  
##   gender = col_integer(),  
##   age = col_integer(),  
##   diabetes = col_integer(),  
##   diabp = col_double(),  
##   sysbp = col_double(),  
##   height = col_integer(),  
##   weight = col_integer(),  
##   smoke = col_integer(),  
##   racebw = col_integer()  
## )
```

# CHD Data

```
dat
```

```
## # A tibble: 3,425 x 17
##       X1    idno choltot incchd   hdl  alcohol   bmi   trig gender   age
##   <int> <int>   <dbl> <int> <int> <dbl> <dbl> <int> <int> <int>
## 1     1     1 3.01e6   224.     0    37     0    36.5   168     1    65
## 2     2     2 3.02e6   259.     0    50    0.06  31.0   136     0    65
## 3     3     3 3.03e6   205.     0    51     4    33.3   114     0    65
## 4     4     4 3.04e6   174.     1    34     0    31.2   332     1    65
## 5     5     5 3.50e6   191.     1    47     0    30.9   139     0    65
## 6     6     6 3.50e6   201.     0    83     0    24.8    91     0    65
## 7     7     7 3.50e6   161.     0    45    8.02  21.7   154     0    65
## 8     8     8 3.50e6   231.     0    33     0    26.6   304     1    65
## 9     9     9 3.50e6   229.     0    48     0    27.2   162     0    65
## 10    10    10 3.50e6   248.     0    47     0    22.6   101     0    65
## # ... with 3,415 more rows, and 7 more variables: diabetes <int>,
## #   diabp <dbl>, sysbp <dbl>, height <int>, weight <int>, smoke <int>,
## #   racebw <int>
```

# CHD Data

- Change gender to female, remove leading column and gender.

```
dat <- dat %>%  
  mutate(female = factor(gender)) %>%  
  select(-gender, -X1)
```



## CHD Data: Replace missing values with means for those covariates.

```
##### Check NAs
sum(is.na(dat))
```

```
## [1] 41
```

```
dat %>% filter(!complete.cases(.))
```

```
## # A tibble: 41 x 16
```

```
##       idno choltot incchd   hdl  alcohol   bmi   trig   age diabetes diabp sysb
##       <int>   <dbl>   <int> <int>   <dbl> <dbl> <int> <int>    <int> <dbl> <dbl>
##  1 3.02e6    222.     0    57  NA    25.8  249    66      1    75     9
##  2 3.50e6    231.     0    42  NA    22.8  182    66      1    81    15
##  3 4.00e6    256.     0    65  NA    32.8  208    66      2    75    12
##  4 5.04e6    247.     0    63  NA    37.1  148    66      1    72    12
##  5 3.04e6    226.     0    53  0.25  NA    105    67      1    77    12
##  6 6.04e6    290.     1    42  NA    22.4  251    67      1    68    12
##  7 3.03e6    213.     0    51  NA    34.9  103    68      2    69    12
##  8 3.04e6    246.     0    47  NA    28.0  212    68      1    50    13
##  9 3.50e6    286.     0    73  NA    35.4  150    68      1    98    16
## 10 3.50e6    249.     0    29  NA    26.6  197    68      2    74    13
## # ... with 31 more rows, and 5 more variables: height <int>, weight <int>,
## #   smoke <int>, racebw <int>, female <fct>
```

# Save cleaned CHD Data

```
write_csv(dat, "dat/chddata_dsi_cleaned.csv")  
dat <- read_csv("dat/chddata_dsi_cleaned.csv")
```

```
## Parsed with column specification:
```

```
## cols(  
##   idno = col_integer(),  
##   choltot = col_double(),  
##   incchd = col_integer(),  
##   hdl = col_integer(),  
##   alcohol = col_double(),  
##   bmi = col_double(),  
##   trig = col_integer(),  
##   age = col_integer(),  
##   diabetes = col_integer(),  
##   diabp = col_double(),  
##   sysbp = col_double(),  
##   height = col_integer(),  
##   weight = col_integer(),  
##   smoke = col_integer(),  
##   racebw = col_integer(),  
##   female = col_integer()  
## )
```

# Using qwraps2 to Summarize Data

```
library(qwraps2)
data_summary <-
  list("Total Chol." =
    list("min" = ~min(choltot),
          "max" = ~max(choltot),
          "mean (sd)" = ~qwraps2::mean_sd(choltot)),
        "Inc. CHD" = list("% (n)" = ~qwraps2::perc_n(incchd)),
        "HDL" =
    list("min" = ~min(hdl),
          "max" = ~max(hdl),
          "mean (sd)" = ~qwraps2::mean_sd(hdl)),
        "BMI" =
    list("min" = ~min(bmi),
          "max" = ~max(bmi),
          "mean (sd)" = ~qwraps2::mean_sd(bmi)))

orig_opt <- options()$qwraps2_markup
options(qwraps2_markup = "markdown")
summary_table(dat, data_summary)
```

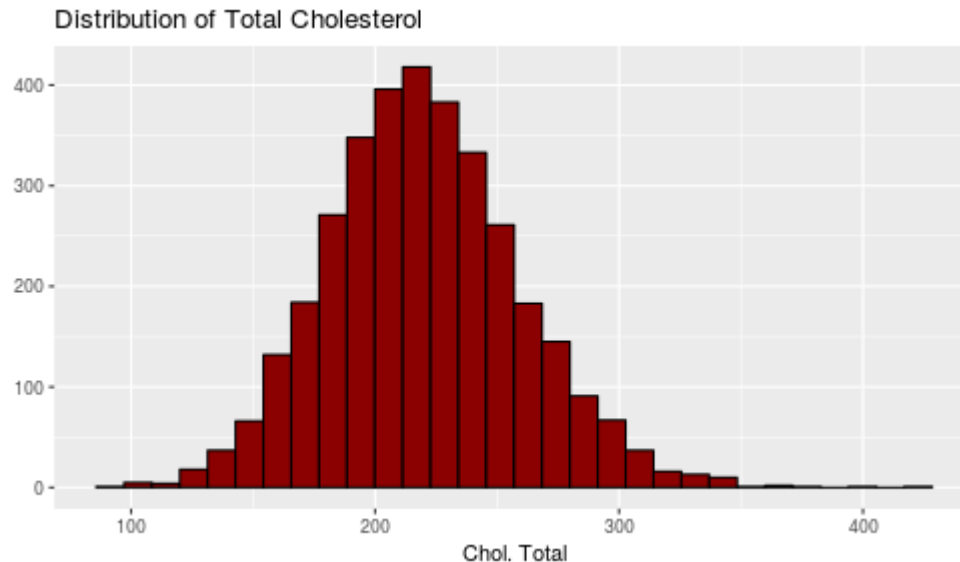
# Using `qwraps2` to Summarize Data

	dat (N = 3425)
<b>Total Chol.</b>	
min	90.59
max	421.87
mean (sd)	219.88 ± 39.48
<b>Inc. CHD</b>	
% (n)	13.66% (n = 3,425)
<b>HDL</b>	
min	15
max	149
mean (sd)	54.90 ± 15.53
<b>BMI</b>	
min	14.65
max	49.41

# Exploratory Plots


```
qplot(choltot, data = dat, fill = I("darkred"), color = I("black"), >
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




# Exploratory Plots

```
qplot(x = factor(incchd), y = choltot, data = dat, geom = "boxplot",
```



# Exploratory Plots

```
qplot(x = factor(incchd), y = log(trig), data = dat, geom = "boxplot")
```



# Bootstrap Estimation

`purrr::map(x, f)` Applies a function  $f$  to each element of a list  $x$

```
a <- 1:10  
purrr::map(a, function(x) x^2)
```

```
## [[1]]  
## [1] 1  
##  
## [[2]]  
## [1] 4  
##  
## [[3]]  
## [1] 9  
##  
## [[4]]  
## [1] 16  
##  
## [[5]]  
## [1] 25  
##  
## [[6]]  
## [1] 36
```





# Bootstrap Estimation

## New Function: `purrr::map`

`purrr::map(x, f)` Applies a function  $f$  to each element of a list  $x$

```
a <- 1:10  
purrr::map_dbl(a, function(x) x^2)
```

```
## [1] 1 4 9 16 25 36 49 64 81 100
```

```
a <- 1:10  
purrr::map_dbl(a, ~ .^2)
```

```
## [1] 1 4 9 16 25 36 49 64 81 100
```

# Bootstrap Estimation

```
boot_mean <- function(d, i) {  
  mean(d[i])  
}
```

```
dat_booted <- dat %>%  
  dplyr::group_by(female) %>%  
  tidyr::nest()  
dat_booted
```

```
## # A tibble: 2 x 2  
##   female data  
##   <int> <list>  
## 1     1 <tibble [1,341 × 15]>  
## 2     0 <tibble [2,084 × 15]>
```







# Bootstrap Estimation

```
dat_booted <- dat_booted %>%  
  dplyr::mutate(booted = purrr::map(.x = data,  
    ~ boot::boot(data = .x$incchd,  
      statistic = boot_mean,  
      R = 5000,  
      stype = "i"))))  
  
dat_booted
```

```
## # A tibble: 2 x 3  
##   female data          booted  
##   <int> <list>          <list>  
## 1     1 <tibble [1,341 × 15]> <S3: boot>  
## 2     0 <tibble [2,084 × 15]> <S3: boot>
```

# Bootstrap Estimation



# Bootstrap Estimation

```
dat_booted <- dat_booted %>%  
  dplyr::mutate(booted_ci = purrr::map(.x = booted,  
                                       ~ boot::boot.ci(.x,  
                                                         conf = 0.95,  
                                                         type = "bca"))  
dat_booted
```

```
## # A tibble: 2 x 4  
##   female data      booted      booted_ci  
##   <int> <list>      <list>      <list>  
## 1       1 <tibble [1,341 × 15]> <S3: boot> <S3: bootci>  
## 2       0 <tibble [2,084 × 15]> <S3: boot> <S3: bootci>
```

# Bootstrap Estimation

```
dat_booted <- dat_booted %>%  
  dplyr::mutate(statistic = purrr::map(.x = booted_ci,  
                                     ~ .x$t0),  
               lower_ci = purrr::map(.x = booted_ci,  
                                     ~ .x$bca[[4]]),  
               upper_ci = purrr::map(.x = booted_ci,  
                                     ~ .x$bca[[5]])) %>%  
  dplyr::select(-data, -booted, -booted_ci) %>%  
  tidyr::unnest()  
dat_booted
```

```
## # A tibble: 2 x 4  
##   female statistic lower_ci upper_ci  
##   <int>      <dbl>    <dbl>    <dbl>  
## 1     1      0.190    0.170    0.210  
## 2     0      0.102    0.0893   0.115
```



# Bootstrap Estimation

```
boot_mean <- function(d, i) {  
  mean(d[i])  
}  
  
dat_booted <- dat %>%  
  dplyr::group_by(female) %>%  
  tidyr::nest() %>%  
  dplyr::mutate(booted = purrr::map(.x = data,  
                                   ~ boot::boot(data = .x$incchd,  
                                                 statistic = boot_mean,  
                                                 R = 5000,  
                                                 stype = "i"))) %>%  
  dplyr::mutate(booted_ci = purrr::map(.x = booted,  
                                       ~ boot::boot.ci(.x,  
                                                       conf = 0.95,  
                                                       type = "bca")),  
  dplyr::mutate(statistic = purrr::map(.x = booted_ci,  
                                       ~ .x$t0),  
    lower_ci = purrr::map(.x = booted_ci,  
                          ~ .x$bca[[4]]),  
    upper_ci = purrr::map(.x = booted_ci,  
                          ~ .x$bca[[5]])) %>%
```

# Bootstrap Estimation

```
knitr::kable(x = dat_booted, digits = 3, col.names = c("Female", "Est
```

95% CI for CHD Incidence by Gener

Female	Est. CHD Incidence	Lower 2.5%	Upper 97.5%
1	0.190	0.169	0.211
0	0.102	0.089	0.116

# Tidy Modeling with broom

```
library(broom)
fit_glm <- glm(incchd ~ ., data = dat, family = "binomial")
tidy(fit_glm)
```

```
## # A tibble: 16 x 5
##   term          estimate std.error statistic    p.value
##   <chr>          <dbl>      <dbl>    <dbl>    <dbl>
## 1 (Intercept) -1.19e+1  4.29      -2.77  0.00553
## 2 idno          4.80e-8  0.00000000456    1.05  0.293
## 3 choltot       3.90e-3  0.00142      2.75  0.00596
## 4 hdl          -6.29e-3  0.00461     -1.36  0.173
## 5 alcohol      -3.16e-2  0.0106     -2.97  0.00299
## 6 bmi           9.70e-2  0.0751      1.29  0.197
## 7 trig          6.76e-4  0.000712      0.950  0.342
## 8 age           3.20e-2  0.00986      3.25  0.00117
## 9 diabetes      2.59e-1  0.0687      3.76  0.000168
## 10 diabp        -6.56e-3  0.00549     -1.20  0.232
## 11 sysbp        1.18e-2  0.00284      4.15  0.0000328
## 12 height       2.87e-2  0.0251      1.14  0.252
## 13 weight      -1.61e-2  0.0125     -1.28  0.199
## 14 smoke        2.00e-1  0.0791      2.53  0.0115
## 15 racebw      -1.86e-2  0.257     -0.0725 0.942
```

# Tidy Modeling with broom

```
knitr::kable(tidy(fit_glm) %>%  
  filter(p.value < 0.05, term != "(Intercept)") %>%  
  select(-statistic) %>%  
  transmute(Source = term, Estimate = estimate, SE = std.error, P = p.value,  
    format = "markdown",  
    digits = 4)
```

Source	Estimate	SE	P
choltot	0.0039	0.0014	0.0060
alcoh	-0.0316	0.0106	0.0030
age	0.0320	0.0099	0.0012
diabetes	0.2587	0.0687	0.0002
sysbp	0.0118	0.0028	0.0000
smoke	0.1999	0.0791	0.0115
female	0.7720	0.1655	0.0000

# Intro to R Markdown

# Intro to R Markdown

- **R Markdown** is an implementation of the *Markdown* document formatting language
- Markdown is a versatile tool that makes it easy to make readable scientific documents in a variety of formats
- R markdown is actively developed and supported by the RStudio team, which means:
  - RStudio has many tools and features to make R Markdown flexible and easy to use
  - New R Markdown features and packages are frequently released

# Intro to R Markdown

- **R For Data Science** on the intent of R Markdown:

*R Markdown files are designed to be used in three ways:*

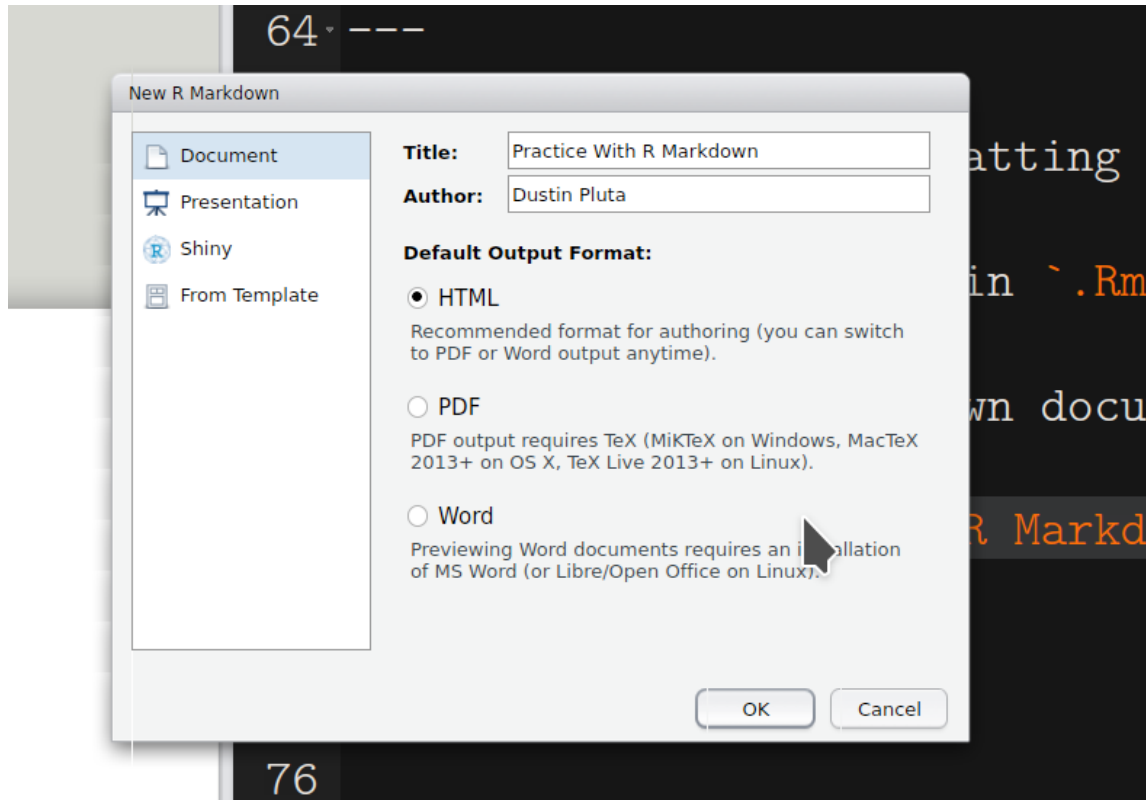
- 1. For communicating to decision makers, who want to focus on the conclusions, not the code behind the analysis.*
- 2. For collaborating with other data scientists (including future you!), who are interested in both your conclusions, and how you reached them (i.e. the code).*
- 3. As an environment in which to do data science, as a modern day lab notebook where you can capture not only what you did, but also what you were thinking.*

# Getting Started with R Markdown

- R Markdown files end in `.Rmd`
- Create a new R markdown document in RStudio:
  - `File > New File > R Markdown...`



# Getting Started with R Markdown



# Getting Started with R Markdown

- The default R Markdown template gives some examples of basic R Markdown features

# Getting Started with R Markdown

# Getting Started with R Markdown

# Getting Started with R Markdown

# Getting Started with R Markdown

# Getting Started with R Markdown

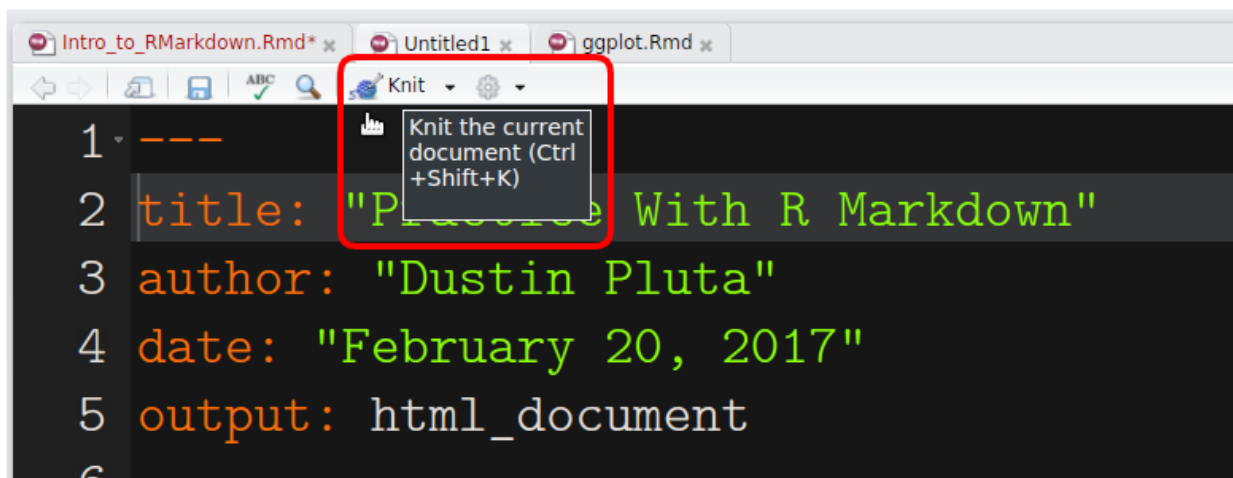
# Getting Started with R Markdown



# Getting Started with R Markdown

# Getting Started with R Markdown

- Compile or "knit" the R Markdown document to the desired format (either html, pdf, or Word document)



# Getting Started with R Markdown

## Practice With R Markdown

*Dustin Pluta*

*February 20, 2017*

### R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:


```
summary(cars)
```

```
##      speed      dist
##  Min.   : 4.0   Min.   : 2.00
##  1st Qu.:12.0   1st Qu.: 26.00
##  Median :15.0   Median : 36.00
##  Mean   :15.4   Mean   : 42.98
##  3rd Qu.:19.0   3rd Qu.: 56.00
##  Max.   :25.0   Max.   :120.00
```

### Including Plots

You can also embed plots, for example:

I




# Getting Started with R Markdown

- Let's modify the plot to include a title, and make the points blue.

```
21
22 ## Including Plots
23
24 You can also embed plots, for example:
25
26 ```{r pressure, echo=FALSE}
27 plot(pressure, main="Plot Title", col="blue", pch=23,
28      bg="blue")
29
30 ```
```

# Getting Started with R Markdown

```
plot(pressure, main = "Plot Title", pch = 23, col = "blue", bg = "blue")
```



# Practice With R Markdown

## R Markdown Cheat Sheet

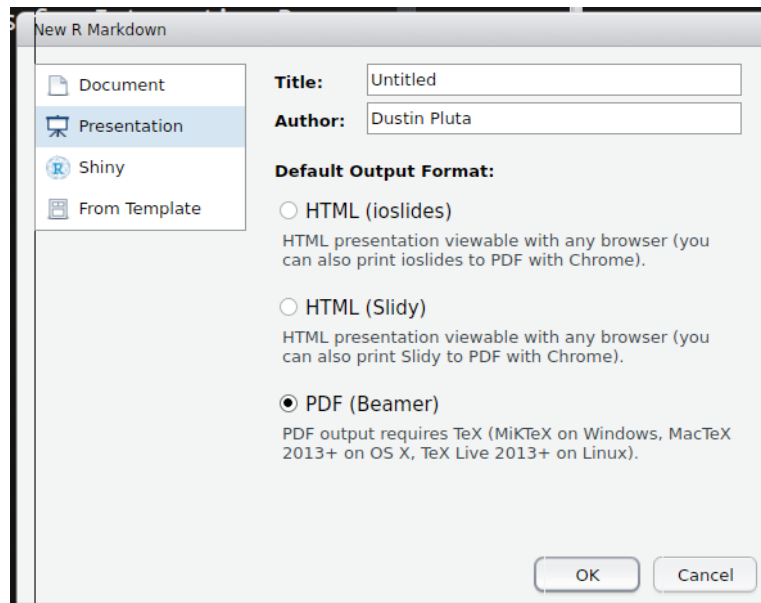
1. R Markdown Basic Example: Shows simple plot with cars data set.
2. R Markdown Exercise Set 1: More examples and some exercises to try on the Iris and IMDB data sets.

# More R Markdown Features

1. Presentations: beamer, ioslides, slidy, xaringan
2. knitr
3. Blogdown
4. Bookdown
5. Interactive Documents

# Presentations

- You can easily create academic presentations using 4 different formats
  - beamer (pdf)
  - ioslides (html)
  - slidy (html)
  - xaringan (html)





# knitr

- R Markdown can make full use of Latex through the `knitr` package
- `knitr` lets you easily display mathematical formulas and other Latex formatting in your Markdown document
- For example, math can be inserted inline like  $\alpha^2 + \beta^2 = \gamma^2$  or in display mode:

$$Y = X\beta + \varepsilon$$

$$\int_{\mathbb{R}} \sum_{i=1}^n \nabla \ell_i d\mu$$

# Blogdown

Making a Website Using Blogdown, Hugo, and GitHub pages/

Example Blogdown Blog: Simply Statistics

# Next Steps

- `RMarkdown_Basic_Example.Rmd`
- `Iris_Example.Rmd`
- `RMarkdown_Exercise_Set1.Rmd`
- `CHD_Data_Analysis.Rmd`

# Some Resources for R

- [dplyr and Data Wrangling Cheat Sheet](#)
- [R Markdown Cheat Sheet](#)
- [Data Carpentry Lessons for R](#)
- [dplyr Tutorial](#)
- [Advanced R](#)
- [R for Data Science](#)
- [Coursera Data Science Specialization](#)